

# CAST: Columnar Agnostic Structural Transformation

A Schema-less Structural Preprocessing Algorithm for Enhancing Compression on Structured and Semi-Structured Data Streams

Andrea Olivari

[github.com/AndreaLVR](https://github.com/AndreaLVR)

January 2026

## Abstract

General-purpose compression algorithms such as LZ77, LZMA, and Zstandard rely on finite dictionary windows and local pattern matching to detect redundancy. While effective for unstructured text, these approaches often fail to fully exploit the long-range structural regularity present in **structured and semi-structured data streams**, including, but not limited to, CSV files, JSON documents, and system logs.

We introduce **CAST** (Columnar Agnostic Structural Transformation), a schema-less structural pre-processing algorithm that infers repetitive layouts directly from the input stream. CAST decomposes each record into a static structural template (*Skeleton*) and a sequence of dynamic values (*Variables*), reorganizing the latter into column-aligned streams prior to compression. This transformation reduces structural entropy and exposes redundancy that is poorly captured by standard compressors operating on row-oriented text.

Experimental evaluation across diverse real-world datasets shows that CAST consistently improves compression density. **Contrary to the traditional trade-off between space and time, CAST reduces end-to-end compression time compared to general-purpose compressors such as LZMA2, Zstandard, and Brotli, when these baselines are configured for their maximum compression efficiency.** By normalizing structural entropy, the algorithm lowers the computational burden on the backend encoder, allowing for simultaneously higher compression ratios and faster throughput. On highly repetitive structured datasets, compression ratios significantly exceed those obtained by standalone compressors, demonstrating that lightweight, schema-free structural normalization can substantially enhance compression density and transmission efficiency without modifying the underlying encoding algorithms.

## 1 Introduction

**Structured and semi-structured data** (logs, IoT telemetry, database dumps, CSV reports) is characterized by rigid and repetitive structural patterns.

Columnar storage formats such as Apache Parquet and ORC exploit this property effectively, but require explicit schemas and are therefore impractical for ad-hoc, semi-structured, or heterogeneous text archives.

Conversely, general-purpose stream compressors such as `xz` (LZMA2) and `zstd` operate without schemas but are constrained by finite dictionary windows and local redundancy detection, often missing long-range structural repetition.

**CAST** proposes a middle ground optimized for **storage efficiency and bandwidth-limited transmission**: a schema-less structural transformation that reorganizes row-oriented text into column-oriented streams inferred dynamically from the data itself. Rather than introducing a new compression algorithm, CAST operates as a pre-processing stage that reshapes the input to better align with the strengths of existing compressors.

This paper details the CAST algorithm: while the transformation is fundamentally backend-agnostic, this study evaluates its effectiveness specifically in combination with **LZMA2**, utilizing a performance-oriented **Reference** Rust implementation<sup>1</sup> (featuring a custom zero-copy parsing strategy).

We test this engine in two configurations: **native library integration** to measure strict algorithmic efficiency (Table 1), and **system-level backend (7-Zip)** to demonstrate real-world throughput. Notably, **this second configuration yields the best overall results**, achieving substantial speedups with negligible compression loss compared to the native integration, as demonstrated in Table 2.

A Python reference implementation is also provided to illustrate the core logic. While functionally analogous to the Rust version, it prioritizes readability via standard Regular Expressions and operates as a single-threaded process in its native configuration, making it distinct from the high-performance Rust engine.

## 2 Methodology

The fundamental premise of CAST is that structured text lines  $L$  can be decomposed into a static template  $S$  and a variable vector  $V$ :

$$L \rightarrow S + V$$

Standard compressors process  $S$  and  $V$  interleaved. CAST processes unique  $S$  sets (Skeletons) once, and  $V$  sets (Variables) as contiguous blocks grouped by column index.

### 2.1 Adaptive Pattern Recognition

The algorithm does not enforce a schema. Instead, it utilizes a **specialized parsing engine** to process lines. To accommodate different data types, it employs an adaptive strategy determined by analyzing the first  $N$  lines (default  $N = 1000$ ) of the input stream.

- **Strict Mode:** Captures quoted strings and explicit numbers. Ideal for **highly structured data with rigid syntax conventions**.
- **Aggressive Mode:** Captures alphanumeric tokens. Ideal for **semi-structured data or free-form text streams**.

**Heuristic thresholds** CAST performs a lightweight statistical analysis on the first  $N$  lines (default  $N = 1000$ ) to select the parsing mode. In our implementation we use the following empirically chosen thresholds: if the ratio of unique skeletons to sample lines exceeds 0.10 we switch from **Strict** to **Aggressive** parsing. To avoid template explosion, the compressor enforces a template budget of  $T_{max} = \alpha \cdot L$  (where  $L$  is number of lines and  $\alpha = 0.25$  in Strict,  $\alpha = 0.40$  in Aggressive); exceeding this budget triggers a safe passthrough (no structural transform).

**Graceful Degradation** The parsing mechanism is designed to handle schema drift without failure. If a line partially matches the **parsing rules**, the unmatched suffix is absorbed into the static skeleton rather than discarded. This ensures that irregular lines simply result in unique templates, allowing the algorithm to naturally degrade towards row-based compression behavior for noisy data, ensuring zero data loss.

---

<sup>1</sup>While this paper validates the *Reference Implementation* (Standard) to establish the algorithmic baseline, a **Work-In-Progress Early Research prototype** of the next-generation architecture (featuring Row Groups for  $O(1)$  Random Access) has been released in the repository’s `rust_random_access_PREVIEW` directory. Preliminary tests on this evolution show promising results, including a 20% ~ 35% **full-file** decompression speedup, though strict validation remains ongoing.

## 2.2 Robustness: The Binary Guard

To prevent data corruption or inefficiency on non-textual files, CAST implements a "Binary Guard". Before processing, a 4KB sample is analyzed. If the density of non-printable control characters (excluding whitespace) exceeds 1%, the input is classified as binary. In this state, CAST enters a "Passthrough" mode, forwarding the raw byte stream directly to the backend compressor with zero structural modification.

## 3 Algorithm Detail

The core transformation logic executes a single-pass processing strategy. To ensure scalability on datasets exceeding physical memory, CAST operates on a streaming block-based architecture, as formally described in Algorithm 1.

---

**Algorithm 1** CAST Streaming Compression Logic

---

```

1: Input: Byte stream  $D$ , Chunk size  $S$  (default:  $S \geq |D|$ )
2: Output: Sequence of compressed blocks  $B_{out}$ 
3: while  $D$  has data do ▷ Streaming loop (main driver)
4:    $C \leftarrow$  read next  $S$  bytes from  $D$ 
   // Step 1: Binary Guard (per-chunk)
5:   if IsLikelyBinary( $C$ ) then
6:     Emit BackendCompress( $C$ ) with flag 255
7:     continue
8:   end if
   // Step 2: Heuristic Analysis (per-chunk)
9:   Strategy  $\leftarrow$  AnalyzeSample( $C$ ) ▷ Strict vs Aggressive
10:  Map  $\leftarrow \{\}$ , Skeletons  $\leftarrow []$ , Columns  $\leftarrow \{\}$ , StreamIDs  $\leftarrow []$ 
    // Step 3: Decomposition (Transformation)
11:  for line in  $C$  do
12:    if ContainsReservedChars(line) then ▷ Collision Fail-Safe
13:      Emit BackendCompress( $C$ ) with flag 255; break
14:    end if
15:     $S, V \leftarrow$  Mask(line, Strategy)
16:    if  $S \notin$  Map then
17:      if templates_exceed_budget() then ▷ Entropy Fail-Safe
18:        Emit BackendCompress( $C$ ) with flag 255; break
19:      end if
20:      Map[ $S$ ]  $\leftarrow$  NewID(), Skeletons.append( $S$ )
21:    end if
22:    Columns[Map[ $S$ ]].append( $V$ ), StreamIDs.append(Map[ $S$ ])
23:  end for
   // Step 4: Finalization (Unified vs Split)
24:  Decision  $\leftarrow$  DecideUnifiedOrSplit(Columns, Skeletons)
25:  if Decision is Unified then
26:    Payload  $\leftarrow$  Pack(Skeletons, StreamIDs, Columns)
27:    Emit BackendCompress(Payload) ▷ e.g., LZMA2, Zstd
28:  else
29:    Emit BackendCompress(Skeletons, StreamIDs, Columns) separately
30:  end if
31: end while

```

---

---

**Algorithm 2** CAST Streaming Reconstruction Logic

---

```
1: Input: Compressed sequence  $B_{in}$ 
2: Output: Reconstructed byte stream  $D_{out}$ 
3: while  $B_{in}$  has data do                                ▷ Process stream chunk-by-chunk
4:   Read Block Header (Lengths, CRC, Flag)
5:   if Flag is 255 (Passthrough) then
6:      $C \leftarrow$  Read compressed payload
7:     Emit BackendDecompress( $C$ )
8:     continue
9:   end if
10:  // Step 1: Backend Inflation
11:   $R, I, V \leftarrow$  BackendDecompress(Payload components)
12:   $Skeletons \leftarrow$  ParseRegistry( $R$ )
13:   $StreamIDs \leftarrow$  ParseIDs( $I$ )
14:  // Step 2: Queue Population ( $O(N)$ )
15:   $Queues \leftarrow$  Initialize variable queues per column from  $V$ 
16:  // Step 3: Row Reconstruction
17:  for  $id$  in  $StreamIDs$  do
18:     $S \leftarrow Skeletons[id]$ 
19:     $Row \leftarrow$  Empty String
20:    for segment in  $S$  do
21:      if segment is Placeholder then
22:         $Var \leftarrow Queues[id].pop\_front()$ 
23:         $Row.append(\text{Unescape}(Var))$ 
24:      else
25:         $Row.append(segment)$ 
26:      end if
27:    end for
28:    Emit  $Row$ 
29:  end for
30: end while
```

---

**Streaming and Scalability** By default, CAST operates in **Solid Mode**, treating the entire input as a single chunk to maximize global redundancy extraction. However, the architecture natively supports **Streaming Chunking** via the parameter  $S$ . This ensures that memory consumption during both compression and decompression remains bounded by  $O(S)$  rather than  $O(D)$ , allowing the engine to process multi-gigabyte logs on commodity hardware without memory exhaustion.

**Unified vs Split decision heuristic** To balance compression density against memory usage, CAST employs a lightweight heuristic primarily based on template cardinality. For datasets with a limited set of unique structures ( $\text{num\_templates} < 256$ ), the algorithm **defaults to Unified Mode** to maximize the shared dictionary context of the backend compressor, **unless a preliminary sampling test detects poor compressibility**. Conversely, when template diversity is high, it strictly enforces **Split Mode** to improve parallelism and reduce the overhead of managing a single monolithic context window.

**Serialization format** Skeletons are concatenated into a registry separated by the Unicode Private Use character U+E001 (selected to avoid collisions with standard text or Latin-1 binaries). Variables are organized per-column and serialized using a row separator (0x00) and a column separator (0x02).

To ensure bitwise reversibility and total binary safety, CAST applies a mandatory **Byte Stuffing** escape scheme regardless of the Unified/Split decision.

The escape byte 0x01 is used to encode reserved control characters (e.g., 0x00 becomes 0x01 0x00). This guarantees a collision-free stream even when processing mixed-encoding logs, binary artifacts, or multi-byte characters, eliminating the alignment failures associated with raw separation.

Conversely, **Split Mode** creates distinct physical streams to maximize parallelism, but adheres to the same escaped binary format to maintain data integrity. The decompressor performs the inverse unescaping prior to reconstruction.

**Collision Safety & Fail-Safe** Although the selected Private Use Area characters (U+E000, U+E001) are disjoint from standard text encodings and unreachable via Latin-1 binary decoding (which maps only up to U+00FF), CAST enforces a deterministic fail-safe mechanism. During the parsing phase, if the input stream is found to naturally contain these reserved markers, the algorithm aborts the transformation and transparently falls back to the backend compressor (Passthrough Mode). This guarantees that the structural reconstruction is never ambiguous, ensuring 100% data integrity even in the theoretical edge case of adversarial inputs.

**Template ID encoding** To minimize metadata overhead, template stream ids are encoded using one of four modes depending on the number of distinct templates:

- **Mode 3:** single template (no ids stored)
- **Mode 2:** 8-bit ids
- **Mode 0:** 16-bit ids
- **Mode 1:** 32-bit ids

The encoder selects the smallest-width representation that can store all template identifiers and records a compact id-mode flag in the header.

**Integrity** For unified (solid) compression CAST stores a compact header with lengths and relies on the underlying compressor integrity check (e.g., LZMA CRC32) and, optionally, an explicit CRC32 verification of the reconstructed payload to ensure bit-perfect lossless round-trip. To guarantee strictly lossless reconstruction (including mixed line-endings like CRLF/LF), the line parsing phase explicitly preserves original terminators.

## 4 Performance Evaluation

To validate the efficacy of CAST, we compiled a heterogeneous corpus of datasets sourced from **Kaggle**, **LogHub**, **Zenodo**, **HuggingFace** and other public repositories. The dataset selection is **intentionally weighted** towards the algorithm’s target domain—structured machine-generated data—to fully explore the optimization potential in relevant scenarios.

However, to define the algorithm’s operational boundaries, we also included a small control group representing **low-redundancy scenarios** (including unstructured text and high-variance structured files). This allows us to transparently verify the hypothesis that CAST’s benefits are strictly dependent on *exploitable* structural redundancy and to quantify the overhead when such redundancy is absent.

To fully evaluate the algorithm’s capabilities, we utilized a **unified high-performance Rust engine** capable of operating in two distinct modes. **While integrated into a single artifact, these modes are architecturally decoupled to prevent benchmark pollution:** this

internal separation ensures that the algorithmic efficiency metrics are derived purely from the transformation logic (Native), strictly isolated from the OS-level piping and binary invocation overhead inherent to the production pipeline (System).

Both configurations feature native **multi-threading** capabilities and configurable memory parameters (specifically `chunk-size` and `dict-size`) to strictly control the runtime footprint (\*\*via block segmentation in compression and streaming I/O in decompression\*\*), preventing memory saturation and enabling the processing of datasets larger than available RAM.

We evaluated this engine in two specific configurations:

- **Native Mode (Reference Configuration):** Selected via the `--mode native` flag, it links directly against native libraries.

**While this configuration fully supports multi-threading and configurable chunking, for the compression ratio analysis it was restricted to single-threaded, monolithic execution to strictly isolate the algorithmic efficiency of the structural transformation without the masking effects of parallelization or context fragmentation.**

- **System Mode (High-Throughput Configuration):** Selected via the `--mode 7zip` flag, it feeds data to the external `7z` executable.

This configuration leverages robust **multi-threading** for industrial scalability, achieving substantially higher speeds with negligible compression loss compared to the Native version.

**Consequently, this benchmark scenario includes many additional large-scale datasets (e.g., 500 MB+) to stress-test the pipeline under heavy load conditions.**

**Backend Flexibility & Tuning.** Crucially, since CAST operates as a transparent pre-processor, the end-to-end performance characteristics are tunable via backend selection.

A user prioritizing restoration speed could pair CAST with a real-time compressor (e.g., Zstd at moderate levels), effectively shifting the trade-off towards "warm" storage requirements. Conversely, for deep archival scenarios where storage density is paramount, a high-ratio backend is preferred. To demonstrate the maximum potential for storage minimization (the primary focus of this work), the following benchmarks evaluate the high-compression **CAST + LZMA2** pipeline.

We benchmarked CAST against three state-of-the-art compression algorithms to provide a comprehensive landscape:

- **LZMA2 (XZ):** Preset 9 (Extreme) with a 128 MB dictionary. To guarantee a strictly fair comparison, we utilize the exact same engine as the CAST backend: the shared native library for the Native Mode tests, and the identical 7-Zip binary/arguments for the System Mode tests.
- **Zstandard (Zstd):** Level 22 (Ultra), representing modern high-performance compression.
- **Brotli:** Quality 11 (Max), widely used for web content optimization.

To ensure a comprehensive assessment, the benchmarks distinguish between three key performance metrics:

1. **Algorithmic Efficiency (Compression Ratio):** The reduction in file size compared to the original raw data, measured using the **Native Mode** to ensure binary-level precision.
2. **Compression Throughput:** We evaluate this primarily using the **System Mode** (7-Zip backend) to leverage an industry-standard, hyper-optimized encoding engine.

This approach demonstrates that CAST achieves **substantially** higher speeds by offloading the heavy lifting to a mature runtime, proving that top-tier performance is attainable without requiring a custom re-implementation of complex multi-threaded encoding logic.

3. **Restoration Latency (Decompression):** Unlike standard algorithms where decompression is a linear byte-stream inflation, CAST requires a *structural reconstruction phase*. We measure this overhead using **both configurations** to differentiate between the intrinsic algorithmic cost (Native Mode) and the end-to-end restoration time including external process management (System Mode).

#### 4.1 Native Mode Benchmarks

**Table 1** details the reduction in file size and processing time across the full dataset corpus using the Rust Native implementation. To visually contextualize these findings without overcrowding the charts, we refer the reader to **Figure 1** for a ranked comparison of algorithmic efficiency on a *representative selection* of these datasets, and **Figure 2** for the corresponding throughput analysis.

Table 1: Native Mode Benchmarks (Single-Threaded Baseline)

Dataset	Original	LZMA2	Zstd	Brothli	CAST (Native)	Ratio
<i>CSV Datasets</i>						
<a href="#">Balance Payments</a>	33.2 MB	501 KB (110s)	698 KB (103s)	592 KB (131s)	<b>245 KB (4.7s)</b>	<b>139x</b>
<a href="#">Migration Stats</a>	29.3 MB	945 KB (56s)	1.12 MB (49s)	1.06 MB (87s)	<b>319 KB (4.6s)</b>	<b>94x</b>
<a href="#">Apple Sitemap</a>	124.2 MB	2.21 MB (178s)	2.51 MB (218s)	2.50 MB (508s)	<b>1.87 MB (34s)</b>	<b>66.4x</b>
<a href="#">DDoS Data</a>	616.8 MB	19.7 MB (1607s)	24.4 MB (1598s)	22.0 MB (1972s)	<b>10.3 MB (417s)</b>	<b>60x</b>
<a href="#">Subnat. Life Tables</a>	16.0 MB	608 KB (13.8s)	713 KB (9.9s)	564 KB (52.5s)	<b>324 KB (3.5s)</b>	<b>50.6x</b>
<a href="#">COVID-19 Surveillance</a>	872.1 MB	25.0 MB (1232s)	27.2 MB (800s)	27.1 MB (3386s)	<b>20.3 MB (559s)</b>	<b>43x</b>
<a href="#">Smart City Sensors</a>	24.2 MB	1.02 MB (54s)	1.14 MB (47s)	1.05 MB (58s)	<b>578 KB (3.6s)</b>	<b>43x</b>
<a href="#">Train/Test Network</a>	29.9 MB	1.05 MB (32.8s)	1.16 MB (26.2s)	1.10 MB (94.2s)	<b>0.89 MB (6.6s)</b>	<b>33.6x</b>
<a href="#">Wireshark</a>	154.4 MB	9.52 MB (457s)	10.8 MB (302s)	10.1 MB (403s)	<b>5.69 MB (167s)</b>	<b>27.1x</b>
<a href="#">RT_IOT2022</a>	54.8 MB	2.53 MB (144s)	2.53 MB (221s)	2.51 MB (47.4s)	<b>1.99 MB (18.5s)</b>	<b>27.5x</b>
<a href="#">NYC Bus Breakdowns</a>	132.9 MB	9.79 MB (165s)	10.4 MB (127s)	10.9 MB (317s)	<b>8.40 MB (93s)</b>	<b>15.8x</b>
<a href="#">NZDep Life Tables</a>	13.1 MB	1.16 MB (8.6s)	1.26 MB (6.5s)	1.14 MB (29s)	<b>881 KB (2.8s)</b>	<b>15.2x</b>
<a href="#">US Stock Prices</a>	224.2 MB	26.9 MB (712s)	29.1 MB (521s)	27.9 MB (495s)	<b>17.4 MB (140s)</b>	<b>13x</b>
<a href="#">Covid Vaccinations</a>	50.8 MB	4.65 MB (48s)	5.09 MB (39s)	4.58 MB (158s)	<b>4.13 MB (21s)</b>	<b>12.3x</b>
<a href="#">HomeC</a>	131.0 MB	14.9 MB (257s)	15.6 MB (195s)	15.6 MB (330s)	<b>11.2 MB (104s)</b>	<b>11.7x</b>
<a href="#">Japan Trade 2020</a>	207.9 MB	24.8 MB (463s)	26.4 MB (355s)	25.2 MB (511s)	<b>18.4 MB (167s)</b>	<b>11.3x</b>
<a href="#">IOT-temp</a>	6.95 MB	788 KB (10.0s)	828 KB (8.6s)	797 KB (16.0s)	<b>728 KB (5.2s)</b>	<b>9.8x</b>
<a href="#">Aus/NZ Fires from space</a>	73.0 MB	9.75 MB (93s)	10.7 MB (67.6s)	10.0 MB (224s)	<b>7.69 MB (52.2s)</b>	<b>9.5x</b>
<a href="#">HAI Security Train</a>	114.2 MB	19.0 MB (123s)	19.6 MB (92s)	19.2 MB (246s)	<b>13.0 MB (67s)</b>	<b>8.8x</b>
<a href="#">IoT Intrusion</a>	197.5 MB	25.1 MB (272s)	25.9 MB (195s)	28.0 MB (521s)	<b>24.0 MB (135s)</b>	<b>8.2x</b>
<a href="#">Nashville Housing</a>	9.9 MB	1.41 MB (8.2s)	1.49 MB (5.9s)	1.41 MB (22s)	<b>1.30 MB (5.3s)</b>	<b>7.6x</b>
<a href="#">Owid Covid</a>	46.7 MB	7.08 MB (55s)	7.49 MB (47s)	6.92 MB (126s)	<b>6.39 MB (26s)</b>	<b>7.3x</b>
<a href="#">Assaults 2015</a>	234 KB	<b>33.9 KB (0.16s)</b>	37.6 KB (0.36s)	34.0 KB (0.41s)	39.6 KB (0.12s)	<b>5.9x</b>
<a href="#">ChatGPT Paraphrases</a>	264.9 MB	<b>40.1 MB (259s)</b>	40.5 MB (200s)	44.4 MB (523s)	44.8 MB (262s)	<b>5.9x</b>
<a href="#">Item Aliases</a>	201.5 MB	40.3 MB (436s)	43.6 MB (301s)	43.4 MB (370s)	<b>40.2 MB (240s)</b>	<b>5x</b>
<a href="#">PaySim Mobile Money</a>	493.5 MB	148.8 MB (756s)	150.8 MB (578s)	154.9 MB (1103s)	<b>130.1 MB (574s)</b>	<b>3.8x</b>
<a href="#">Dielectron Collision</a>	14.7 MB	5.71 MB (14.2s)	6.08 MB ( <b>12.0s</b> )	5.88 MB (33s)	<b>5.26 MB (12.6s)</b>	<b>2.8x</b>
<i>Logs and SQL Datasets</i>						
<a href="#">OpenSSH Logs (LogHub)</a>	70.02 MB	2.23 MB (126.8s)	2.84 MB (112.0s)	2.61 MB (187.3s)	<b>1.01 MB (18.6s)</b>	<b>69.3x</b>
<a href="#">PostgreSQL JSON Logs (Zenodo)</a>	392.8 MB	9.29 MB (1674s)	10.5 MB (1434s)	10.4 MB (1560s)	<b>6.20 MB (170s)</b>	<b>63.3x</b>
<a href="#">Spider NLP Train</a>	23.2 MB	482 KB (11.4s)	483 KB (15.6s)	509 KB (24.5s)	<b>406 KB (3.4s)</b>	<b>57.1x</b>
<a href="#">BGL (LogHub)</a>	708.8 MB	26.6 MB (1330s)	31.2 MB (942s)	33.0 MB (2051s)	<b>19.7 MB (242s)</b>	<b>36.0x</b>
<a href="#">Sakila DB</a>	8.8 MB	427 KB (26s)	502 KB (21s)	466 KB (28s)	<b>298 KB (2.2s)</b>	<b>29.5x</b>
<a href="#">Weblog Sample</a>	67.6 MB	2.65 MB (67s)	2.88 MB (60s)	3.07 MB (204s)	<b>2.53 MB (34s)</b>	<b>26.7x</b>
<a href="#">Apache Server Logs</a>	242.0 MB	12.9 MB (306s)	13.3 MB (217s)	14.2 MB (679s)	<b>10.3 MB (131s)</b>	<b>23.5x</b>
<a href="#">Cross Reference</a>	12.0 MB	1.15 MB (15.7s)	1.28 MB (12.1s)	1.34 MB (26s)	<b>1.08 MB (9.1s)</b>	<b>11.1x</b>
<a href="#">World Cities SQL</a>	20.3 MB	<b>2.47 MB (22.5s)</b>	2.78 MB (17.0s)	2.54 MB (48.2s)	2.48 MB ( <b>10.4s</b> )	<b>8.2x</b>
<a href="#">Bible MySQL</a>	32.3 MB	<b>4.30 MB (39s)</b>	4.45 MB ( <b>32s</b> )	6.42 MB (69s)	4.30 MB (42s)	<b>7.5x</b>
<a href="#">mssql YLT Table SQL</a>	12.8 MB	<b>1.02 MB (30.9s)</b>	1.15 MB ( <b>27.4s</b> )	1.11 MB (31.8s)	1.02 MB (31.3s)	<b>12.5x</b>
<a href="#">Audit Dump (SQL)</a>	64.6 MB	12.0 MB (135s)	12.6 MB (109s)	12.1 MB (146s)	<b>10.2 MB (35s)</b>	<b>6.3x</b>
<a href="#">Pixel 4XL GNSS Log</a>	130.5 MB	32.1 MB (207s)	33.3 MB ( <b>150s</b> )	31.5 MB (297s)	<b>25.0 MB (219s)</b>	<b>5.2x</b>
<i>JSON &amp; XML Datasets</i>						
<a href="#">CERN OpenData Slim</a>	211.9 MB	16.1 MB (393s)	17.0 MB (362s)	16.4 MB (510s)	<b>15.2 MB (327s)</b>	<b>14x</b>
<a href="#">Yelp Business</a>	118.9 MB	<b>9.87 MB (123s)</b>	10.2 MB (102s)	11.2 MB (293s)	10.4 MB ( <b>73s</b> )	<b>11.4x</b>
<a href="#">Brazil Geo</a>	14.5 MB	1.55 MB (8.2s)	1.64 MB ( <b>6.1s</b> )	1.67 MB (27s)	<b>1.53 MB (9.8s)</b>	<b>9.5x</b>
<a href="#">Wikidata Fanout</a>	262.3 MB	33.0 MB (284s)	38.2 MB (205s)	34.8 MB (533s)	<b>28.5 MB (181s)</b>	<b>9.2x</b>
<a href="#">Pagerank</a>	121.9 MB	14.8 MB ( <b>127s</b> )	15.2 MB (136s)	16.5 MB (289s)	<b>14.7 MB (129s)</b>	<b>8.3x</b>
<a href="#">Yelp Tips</a>	180.6 MB	34.2 MB (188s)	34.7 MB (132s)	44.1 MB (376s)	<b>30.6 MB (107s)</b>	<b>5.9x</b>
<a href="#">Yelp Checkin</a>	287.0 MB	54.1 MB (500s)	61.9 MB ( <b>386s</b> )	59.7 MB (866s)	<b>53.6 MB (422s)</b>	<b>5.4x</b>
<a href="#">Gandhi Works</a>	100.6 MB	20.6 MB (106s)	20.9 MB ( <b>84s</b> )	22.5 MB (217s)	<b>20.3 MB (92s)</b>	<b>5x</b>
<a href="#">Wiki 00c2bfc7-.json</a>	41.2 MB	10.26 MB (51s)	10.50 MB ( <b>41s</b> )	10.80 MB (107s)	<b>10.24 MB (49s)</b>	<b>4x</b>
<a href="#">GloVe Emb.</a>	193.4 MB	57.7 MB (400s)	57.9 MB ( <b>231s</b> )	60.3 MB (501s)	<b>57.1 MB (401s)</b>	<b>3.4x</b>
<i>Text Datasets</i>						
<a href="#">XDados</a>	4.4 MB	535 KB (3.1s)	597 KB ( <b>2.2s</b> )	536 KB (9.7s)	<b>420 KB (3.4s)</b>	<b>10.5x</b>
<a href="#">Tatoeba French</a>	26.9 MB	<b>3.69 MB (24.5s)</b>	3.87 MB ( <b>20.5s</b> )	3.81 MB (51s)	5.14 MB (25s)	<b>5.2x</b>
<a href="#">Election Day Tweets</a>	35.9 MB	<b>10.1 MB (31s)</b>	10.2 MB ( <b>24s</b> )	10.4 MB (74s)	12.1 MB (38s)	<b>3x</b>

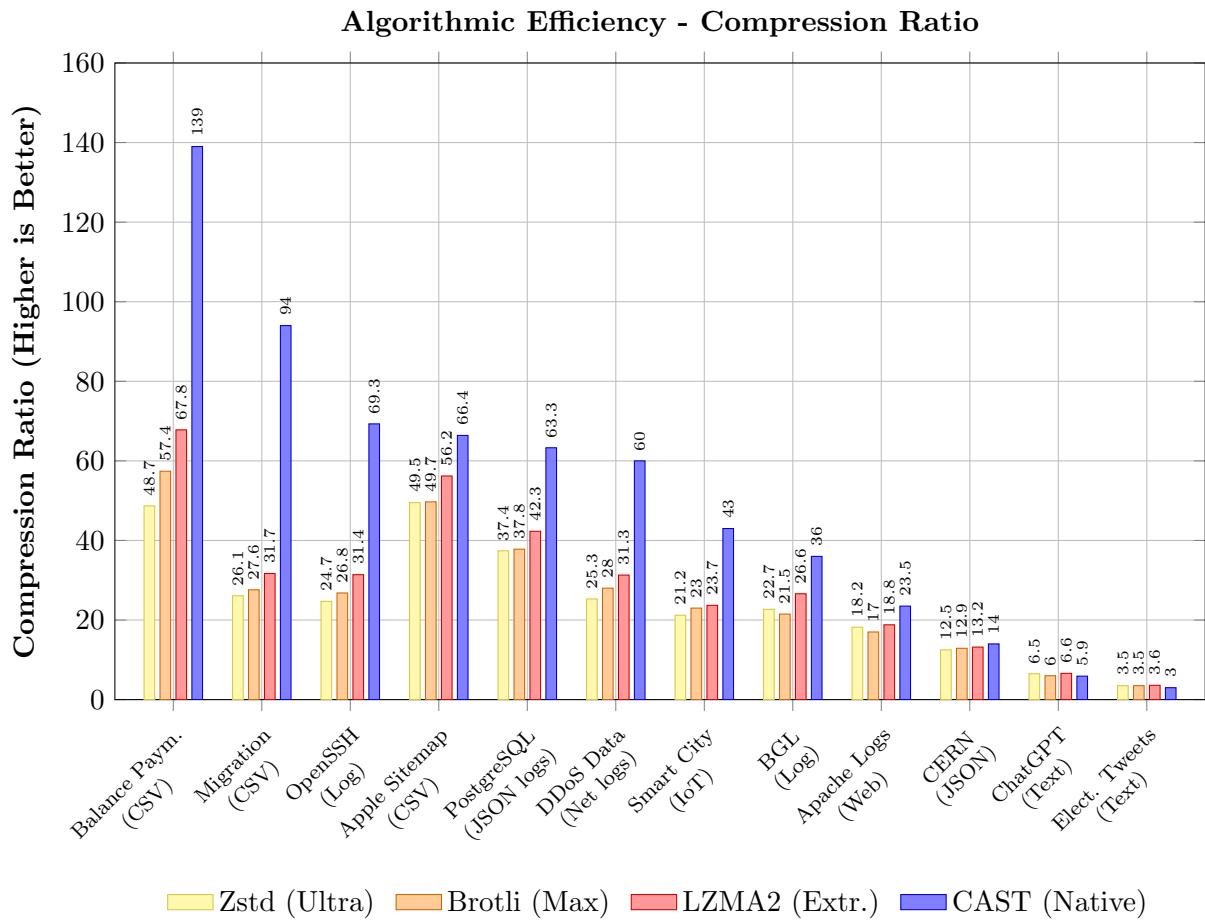


Figure 1: **Algorithmic Efficiency Gap:** Comparison of compression ratios on a comprehensive selection of datasets (sorted by CAST performance). CAST (Blue) shows dominant performance on structured data, doubling or tripling competitor density. The inclusion of unstructured/high entropy text (e.g., *Elect. Tweets*) demonstrates the natural boundary where structural pre-processing becomes neutral or slightly disadvantageous.

### Compression Throughput Comparison (Native Mode - Single Thread)

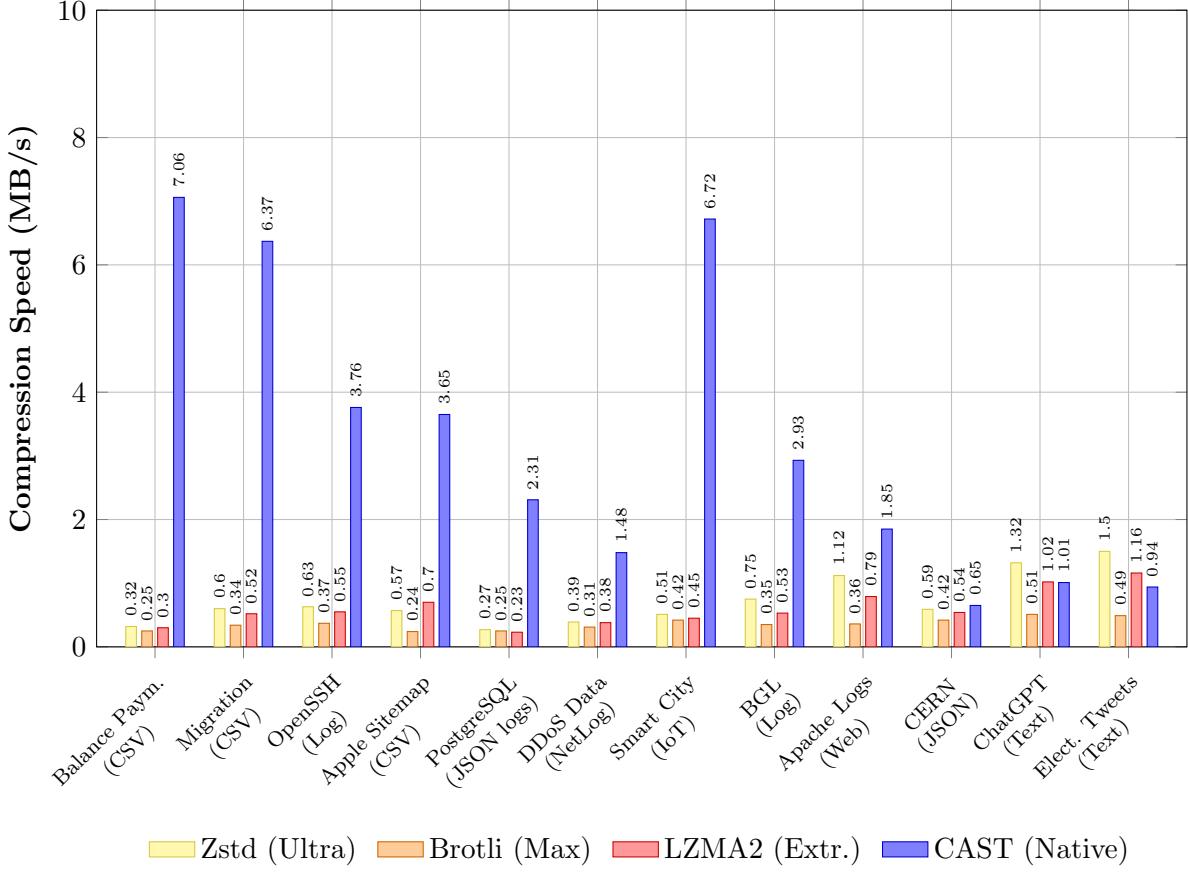


Figure 2: **Throughput Analysis (Native Mode):** Even in single-threaded mode, CAST (Blue) achieves significantly higher throughput on CSVs and structured/semi-structured logs due to the reduced input size passed to the backend. On unstructured/high entropy text (*ChatGPT*, *Tweets*), the parsing overhead aligns performance with standard compressors.

## 4.2 System Mode Benchmarks

A common drawback of pre-processing is added latency. However, Table 2 demonstrates that offloading the compression workload to the highly optimized 7-Zip backend not only eliminates the overhead observed in the Native implementation, but actually achieves higher throughput than applying 7-Zip directly to the raw file.

**Controlled Environment & Analysis:** To ensure a strictly fair comparison, the control group ("LZMA2 (7-Zip)") employs the exact same compression binary and threading configuration used by the CAST backend.

Consequently, the observed speedup is attributable solely to the entropy reduction achieved by the structural transformation.

Although CAST adds a parsing step, the resulting column-aligned streams (e.g., continuous integers or timestamps) are significantly less computationally expensive for the LZMA encoder to process than the raw chaotic text, resulting in a net reduction of total processing time.

Crucially, the benchmarks honestly report the performance regression on unstructured text (e.g., ChatGPT, Tweets), where the parsing overhead is not offset by compression gains. This confirms the hypothesis that CAST's advantage is strictly structural.

To visualize this phenomenon, **Figure 3** correlates the processing throughput with the compression ratio. The data reveals a clear trend: as the structural density increases (higher ratio), the backend encoding speed accelerates proportionally, confirming that CAST acts as a throughput multiplier for structured data.

Table 2: System Mode Benchmarks

Dataset	Original	LZMA2 (7-Zip)	CAST (Native)	CAST (System)	Compression Ratio	Speedup (vs LZMA2)
<i>CSV Datasets</i>						
Balance Payments	33.1 MB	834 KB (2.0s)	<b>245 KB (4.7s)</b>	255 KB ( <b>0.75s</b> )	138x	<b>2.69x</b>
Migration Stats	29.2 MB	1.38 MB (4.6s)	<b>319 KB (4.6s)</b>	343 KB ( <b>1.16s</b> )	93.7x	<b>4.00x</b>
Sitemap Apple	124.2 MB	2.69 MB (9.25s)	<b>1.87 MB (34s)</b>	1.99 MB ( <b>6.51s</b> )	66.4x	<b>1.42x</b>
Sowiport Rec Logs	2.30 GB	51.0 MB (189s)	-	<b>39.3 MB (139s)</b>	60x	<b>1.36x</b>
DDoS Data	616.8 MB	20.4 MB (81s)	<b>10.3 MB (417s)</b>	10.9 MB ( <b>37s</b> )	60x	<b>2.17x</b>
Subnat. Life Tables	16.0 MB	824 KB (2.6s)	<b>324 KB (3.5s)</b>	344 KB ( <b>0.8s</b> )	<b>50.5x</b>	<b>3.27x</b>
Smart City Sensors	23.04 MB	1.23 MB (5.8s)	578 KB (3.6s)	<b>0.57 MB (1s)</b>	41.9x	<b>5.88x</b>
COVID-19 Surveillance	872.1 MB	32.8 MB (63s)	-	<b>22.3 MB (50s)</b>	39x	<b>1.26x</b>
Train/Test Network	28.5 MB	1.29 MB (3.9s)	<b>0.89 MB (6.6s)</b>	<b>0.89 MB (1.8s)</b>	32x	<b>2.14x</b>
RT-IOT2022	54.8 MB	2.56 MB (8.7s)	<b>1.99 MB (18.5s)</b>	2.01 MB ( <b>5.9s</b> )	<b>27.5x</b>	<b>1.47x</b>
Wireshark P3	154.4 MB	10.6 MB (47.7s)	<b>5.69 MB (167s)</b>	6.94 MB ( <b>32s</b> )	27x	<b>1.49x</b>
Japan Trade 2018	668.3 MB	56.6 MB (105s)	-	<b>25.9 MB (78s)</b>	<b>25.8x</b>	<b>1.33x</b>
SQL Injection	1004.5 MB	<b>60.1 MB (97s)</b>	-	62.1 MB (182s)	16x	0.54x
NYC Bus Breakdowns	126.71 MB	10.55 MB (40.2s)	8.40 MB (93s)	<b>8.24 MB (24.1s)</b>	<b>15.3x</b>	<b>1.67x</b>
NZDep Life Tables	13.0 MB	1.20 MB (2.7s)	<b>881 KB (2.8s)</b>	882 KB ( <b>1.08s</b> )	15x	<b>2.53x</b>
H1B Data	756.6 MB	<b>45.5 MB (106s)</b>	-	53.7 MB ( <b>95s</b> )	14x	<b>1.12x</b>
US Stock Prices	213.82 MB	27.18 MB (94.1s)	<b>17.4 MB (140s)</b>	16.69 MB ( <b>48s</b> )	<b>12.8x</b>	<b>1.96x</b>
HomeC	131.0 MB	15.4 MB (54.6s)	<b>11.2 MB (104s)</b>	11.7 MB ( <b>35.7s</b> )	<b>11.7x</b>	<b>1.53x</b>
Japan Trade 2020	207.9 MB	25.3 MB (89s)	<b>18.4 MB (167s)</b>	19.0 MB ( <b>53.2s</b> )	<b>11.3x</b>	<b>1.68x</b>
IOT Temp	6.9 MB	787 KB (1.45s)	728 KB (5.2s)	<b>724 KB (0.94s)</b>	9.7x	<b>1.54x</b>
Aus/NZ Fires from space	73.0 MB	9.74 MB (30.2s)	<b>7.69 MB (52.2s)</b>	7.76 MB ( <b>16.1s</b> )	9.5x	<b>1.87x</b>
HAI Security Train	108.91 MB	18.27 MB (53.1s)	13.0 MB (67s)	<b>12.4 MB (29.8s)</b>	8.8x	<b>1.78x</b>
IoT Intrusion	197.5 MB	28.2 MB (99.7s)	<b>24.0 MB (135s)</b>	24.2 MB ( <b>63.5s</b> )	8.2x	<b>1.57x</b>
Nashville Housing	9.9 MB	1.42 MB (2.4s)	1.30 MB (5.3s)	<b>1.28 MB (1.7s)</b>	7.7x	<b>1.36x</b>
OWID Covid	46.7 MB	7.20 MB (15.7s)	6.39 MB (26s)	<b>6.36 MB (13.2s)</b>	7.3x	<b>1.19x</b>
NASA Global Fire Data	502.2 MB	86.7 MB (177.5s)	-	<b>72.7 MB (140.8s)</b>	6.9x	<b>1.26x</b>
Assaults 2015	234 KB	<b>34.4 KB (0.06s)</b>	39.6 KB (0.12s)	39.9 KB (0.07s)	5.9x	0.86x
ChatGPT Paraphrases	252.6 MB	<b>38.3 MB (142s)</b>	44.8 MB (262s)	42.6 MB (151s)	<b>5.9x</b>	0.94x
Satellites Solar Wind	873.9 MB	211.5 MB (317.5s)	-	<b>161.5 MB (214s)</b>	5.4x	<b>1.48x</b>
Item Aliases	201.5 MB	40.6 MB ( <b>83.7s</b> )	40.2 MB (240s)	40.1 MB (91.1s)	5x	0.92x
PaySim Mobile Money	470.6 MB	143.27 MB (283s)	130.1 MB (574s)	<b>125.6 MB (264s)</b>	3.7x	<b>1.07x</b>
Synthetic Financial Log	470.7 MB	143.3 MB (362s)	-	<b>125.7 MB (263s)</b>	3.7x	<b>1.38x</b>
Dielectron Collision	14.06 MB	5.44 MB (8.46s)	5.26 MB (12.6s)	<b>4.99 MB (8.2s)</b>	2.8x	<b>1.03x</b>
<i>Logs &amp; SQL Datasets</i>						
OpenSSH Logs (LogHub)	70.02 MB	3.43 MB (11.2s)	-	<b>1.01 MB (2.9s)</b>	<b>69.3x</b>	<b>3.84x</b>
PostgreSQL JSON Logs (Zenodo)	392.83 MB	12.5 MB (36s)	<b>6.2 MB (165s)</b>	6.67 MB ( <b>25s</b> )	<b>63.3x</b>	<b>1.44x</b>
IOTA logs 2.21 (merge)	1.48 GB	73.1 MB (83.6s)	-	<b>40.8 MB (80.2s)</b>	<b>37.1x</b>	<b>1.04x</b>
BGL (LogHub)	708.76 MB	31.49 MB (74.3s)	-	<b>20.1 MB (46.5s)</b>	<b>35.6x</b>	<b>1.60x</b>
Web Server Logs Labeled	2.66 GB	103.3 MB (245s)	-	<b>86.0 MB (181s)</b>	<b>31.6x</b>	<b>1.36x</b>
Sakila Insert	8.8 MB	492 KB (0.97s)	298 KB (2.2s)	<b>297 KB (0.53s)</b>	30.3x	<b>1.83x</b>
HDFS v1 (LogHub) *	1.47 GB	91.1 MB (194s)	-	<b>49.8 MB (230s)</b>	<b>30.2x</b>	0.84x
Apache Server Logs	242.0 MB	15.7 MB (52.6s)	<b>10.3 MB (131s)</b>	11.9 MB ( <b>33.2s</b> )	<b>23.5x</b>	<b>1.58x</b>
Cross Reference	11.5 MB	1.11 MB (3.72s)	1.08 MB (9.1s)	<b>1.04 MB (2.4s)</b>	11x	<b>1.52x</b>
Bible MySQL	30.8 MB	<b>4.16 MB (15s)</b>	4.30 MB (42s)	<b>4.16 MB (16s)</b>	7.4x	0.94x
Audit Dump (SQL)	64.6 MB	12.4 MB (27.1s)	10.2 MB (35s)	<b>10.0 MB (14.3s)</b>	6.5x	<b>1.90x</b>
<i>JSON Datasets</i>						
CERN OpenData Slim	202.12 MB	16.66 MB (49.76s)	<b>15.2 MB (327s)</b>	15.42 MB ( <b>45.6s</b> )	<b>13.3x</b>	<b>1.09x</b>
Yelp Business	118.9 MB	11.1 MB (32.5s)	<b>10.4 MB (73s)</b>	10.9 MB ( <b>23.6s</b> )	<b>11.4x</b>	<b>1.38x</b>
Brazil Geo	14.5 MB	1.55 MB ( <b>2.5s</b> )	1.53 MB (9.8s)	<b>1.5 MB (3s)</b>	<b>9.6x</b>	0.83x
Wikidata Fanout	262.3 MB	33.4 MB (139s)	28.5 MB (181s)	<b>27.8 MB (119s)</b>	9.4x	<b>1.17x</b>
Pagerank	121.9 MB	15.8 MB ( <b>45.4s</b> )	<b>14.7 MB (129s)</b>	15.7 MB (48.6s)	<b>8.3x</b>	0.93x
Yelp Tips	180.6 MB	35.0 MB (79.3s)	30.6 MB (107s)	<b>30.4 MB (53s)</b>	<b>5.9x</b>	<b>1.50x</b>
Yelp Checkin	287.0 MB	55.0 MB ( <b>157s</b> )	<b>53.6 MB (422s)</b>	54.2 MB (169s)	<b>5.3x</b>	0.93x
Gandhi Works	100.6 MB	20.8 MB (55.4s)	<b>20.3 MB (92s)</b>	<b>20.3 MB (55.2s)</b>	5x	<b>1.00x</b>
Wiki 00c2bfc7-.json	41.2 MB	10.3 MB ( <b>17.7s</b> )	<b>10.2 MB (49s)</b>	<b>10.2 MB (18.6s)</b>	4x	0.95x
Glove Emb.	193.4 MB	58.1 MB ( <b>179s</b> )	<b>57.1 MB (401s)</b>	57.8 MB (195s)	<b>3.4x</b>	0.92x
<i>Text Datasets</i>						
Russian Dictionary	269.2 MB	<b>13.9 MB (109s)</b>	-	13.9 MB (119s)	<b>19.37x</b>	0.91x
Xdados	4.4 MB	533 KB (1.1s)	<b>420 KB (3.4s)</b>	433 KB (1.26s)	<b>10.7x</b>	0.87x
Tatoeba French	26.9 MB	<b>3.74 MB (8.1s)</b>	5.14 MB (25s)	5.02 MB (9.3s)	<b>5.36x</b>	0.87x
Pixel 4XL GNSS Log	124.43 MB	30.87 MB (90.6s)	25.0 MB (219s)	<b>24.3 MB (76.1s)</b>	<b>5.1x</b>	<b>1.19x</b>
Election Day Tweets	34.26 MB	<b>9.64 MB (13.4s)</b>	12.1 MB (38s)	11.27 MB (17.3s)	3x	0.78x

\* Configured with a chunk-size of 300 MB (vs default no chunking). Contrarily to typical behavior, this segmentation yielded superior compression compared to a monolithic (single-chunk) strategy, demonstrating that isolating local entropy in heterogeneous streams can significantly enhance backend efficiency.




Figure 3: **System Mode Analysis.** This combined view illustrates the relationship between structural compression and throughput. High structural redundancy (peaks in the Green Line) generally acts as a catalyst for performance, enabling CAST (Blue Bars) to significantly outpace the LZMA2 baseline (Red Bars). While system-level factors (e.g., I/O on large datasets) prevent strictly linear scaling, the data confirms that superior compression ratios are the primary driver of accelerated processing speed.

### 4.3 Decompression and Reconstruction Analysis

CAST involves an additional *reconstruction phase*: the algorithm must re-interleave the decompressed columnar data back into the original row-based textual format. Table 3 compares the restoration time between the Rust Native implementation and the Rust implementation using the 7-Zip backend.

#### Analysis.

- **Reconstruction Throughput:** The reconstruction phase introduces computational overhead compared to raw stream inflation. However, as shown in **Figure 4**, throughput remains high on structured inputs. CAST achieves **50–115 MB/s** on **highly structured datasets** (e.g., sensors, logs, CSVs), validating its suitability for **archival restoration** and **batch processing**, where aggregate throughput is sufficient to saturate standard network links.
- **Native vs. External Backend:** The System Mode (Rust + 7-Zip) generally outperforms the Native version on larger files due to 7-Zip’s highly optimized multi-threaded decoding

engine. However, the Native implementation is competitive and occasionally faster on smaller datasets, likely due to the absence of process-spawning overhead.

Table 3: Decompression Time Comparison: Native vs System

Dataset	Original Size	Compressed Size ( $\approx$ )	Native Mode	System Mode	Best Speed (MB/s)
Smart City Sensors	23.04 MB	0.57 MB	0.22s	<b>0.20s</b>	<b>115.2</b>
HDFS v1 (LogHub)	1.47 GB	50 MB	14.20s	<b>14.10s</b>	<b>106.7</b>
BGL (LogHub)	708.8 MB	20 MB	9.36s	<b>8.20s</b>	<b>86.4</b>
Recipes JSON	85.2 MB	6 MB	1.34s	<b>1.04s</b>	<b>81.9</b>
PostgreSQL JSON Logs (Zenodo)	392.8 MB	6.5 MB	<b>5.35s</b>	5.62s	<b>73.4</b>
US Stock Prices	213.82 MB	16.7 MB	3.90s	<b>3.29s</b>	<b>65.0</b>
CERN OpenData Slim	202.12 MB	15 MB	4.10s	<b>3.21s</b>	<b>63.0</b>
IOTA logs 2.21 (merge)	1.48 GB	41 MB	34.60s	<b>26.17s</b>	<b>57.9</b>
HAI Security Train	108.91 MB	12 MB	2.59s	<b>1.93s</b>	<b>56.4</b>
ChatGPT Paraphrases	252.6 MB	43 MB	5.67s	<b>4.70s</b>	<b>53.7</b>
Pixel 4XL GNSS Log	124.43 MB	24 MB	2.92s	<b>2.33s</b>	<b>53.4</b>
PaySim Mobile Money	470.6 MB	125 MB	13.76s	<b>9.59s</b>	<b>49.1</b>
NYC Bus Breakdowns	126.71 MB	8 MB	<b>3.14s</b>	3.25s	<b>40.4</b>

**Theoretical Complexity:** Unlike iterative compression which requires costly pattern matching ( $O(N \cdot W)$  where  $W$  is window size), the reconstruction phase is strictly linear  $O(N)$ . The decoder operates via buffered streaming and direct memory copying of pre-calculated offsets. Consequently, the structural restoration imposes **negligible computational overhead**, ensuring that the total decompression time remains dominated by the underlying backend decoding rather than the reconstruction logic.




Figure 4: **Decompression Performance Landscape.** The chart displays the peak decompression speed achieved by CAST. Labels above bars indicate the original dataset size. Notably, CAST maintains high throughput ( $> 100$  MB/s) on massive datasets like HDFS (1.5 GB). Performance correlates with *internal redundancy* rather than just file format: highly repetitive streams (left) are decoded rapidly, whereas high-entropy datasets—whether unstructured text or complex CSVs (e.g., *NYC Bus*)—yield lower throughput.

## 5 Limitations & Constraints

- 1. Binary & High Entropy Data:** As evidenced by benchmarks, CAST provides no benefit for binary/high entropy data. The Binary Guard correctly identifies these files, resulting in performance identical to standard LZMA2 (Equal size/time), but no gain.
- 2. Vector Data:** On datasets consisting primarily of high-variance floating point numbers (e.g., GloVe Embeddings), the structural overhead matches the compression gain, resulting in a ratio gain of only 1.00x.
- 3. Small Files:** For files under 1MB, the overhead of the CAST header and dictionary structure may result in a slightly larger file size compared to raw compression.
- 4. Maturity & Performance Profile:** Significant engineering effort has been dedicated to ensuring the robustness and stability of the provided implementations, incorporating strict memory safeguards and error handling strategies. However, as a recently developed research prototype, it naturally lacks the extensive, multi-year validation and security auditing characteristic of legacy standard utilities like `xz` or `zstd`. Nevertheless, empirical testing suggests a high degree of reliability and efficiency; in particular, the **System Mode** pipeline demonstrates exceptional potential, delivering throughput that frequently exceeds

the standalone 7-Zip baseline, validating the practical viability of the approach. Future development is committed to addressing emerging edge cases and further hardening the codebase, ensuring a continuous evolution towards full production readiness.

5. **Backend Agnosticism & Scalability:** CAST functions as a generic pre-processor compatible with any stream compressor. This architecture allows users to optimize the pipeline: pairing CAST with high-throughput backends (e.g., Zstd) **directly improves end-to-end restoration speed** by minimizing the raw decompression phase. While the structural reconstruction step imposes a fixed computational baseline, the overall system scales effectively with the backend’s performance, allowing users to select the optimal trade-off between compression ratio (LZMA2) and access speed (Zstd).

## 6 Related Work

CAST intersects with several areas in data representation and compression. Here we position our contribution relative to prior art, specifically distinguishing between analytic-focused storage and archival compression.

**Columnar storage** Formats such as Apache Parquet and ORC store data column-wise to optimize analytic workloads (OLAP). However, they generally require explicit schemas definition upfront and are widely adopted for structured data lakes rather than raw log archival. CAST aims to bring similar entropy reduction benefits to schema-less, textual inputs without enforcing strict typing or migration to a new storage format.

**Log-specific and Structural Compressors** State-of-the-art tools like CLP, Crystal, and frameworks like OpenZL operate on principles similar to CAST: they perform structural decomposition (separating static templates from dynamic variables) without requiring an upfront manually defined schema. However, a key distinction lies in the design goal.

Systems like CLP and Crystal are **tailored specifically for log data** and often prioritize *searchability* and granular random access (e.g., query execution directly on compressed blocks). In contrast, CAST is designed primarily for *cold storage* and maximum archival density. It acts as a pure transformation layer to feed strong backend compressors (like LZMA2), prioritizing reduction ratio over immediate query latency.

While CAST currently utilizes chunking for memory safety, this block-based architecture shares the potential for future indexed random access implementations, similar to the aforementioned systems.

**Preconditioning and Transform Filters** Transforms such as Delta encoding, Burrows–Wheeler Transform (BWT), and dictionary pre-filters are commonly used as preprocessing steps to aid compression. CAST fits into this category but targets *structural* redundancy (syntax and hierarchy) rather than numeric locality or symbol re-ordering, making it orthogonal and complementary to these filters.

**Summary** CAST does not seek to replace analytic formats or search-optimized log compressors. Instead, it serves as a lightweight, reversible structural normalization layer. It is particularly effective for long-term archival scenarios where the primary metric is storage efficiency, bridging the gap between raw text compression and specialized structural encoding.

## 7 Reproducibility

Source code and benchmarking scripts are available in the project repository.

To facilitate immediate adoption, the tool requires minimal setup and features a streamlined command-line interface (CLI) designed for intuitive usage—analogous to standard utilities like 7-Zip.

To ensure reproducibility without distributing large third-party files, the datasets used are all publicly available, with their specific sources cited directly in this paper and listed in the repository documentation.

Benchmarks were executed on a commodity workstation running **Windows 10 Pro for Workstations** (Build 19045), equipped with a 6th-generation **Intel Core i7 CPU** (Skylake architecture, 3.40 GHz) and 16 GB of physical RAM. Experiments were conducted using native Rust builds (`rustc 1.92.0`), Python 3.10 and 7-Zip 23.01.

Exact invocation commands are documented in the repository.

## 8 Conclusion

This work demonstrates that structural pre-processing can significantly enhance the effectiveness of general-purpose compression algorithms when applied to structured text data. By dynamically inferring and separating static templates from variable fields, CAST reshapes row-oriented text into column-aligned streams that better expose redundancy to standard compressors.

While CAST does not provide benefits for binary or high-entropy data, empirical results show substantial improvements in both compression ratio and throughput for logs, CSV files, and semi-structured datasets. These findings indicate that schema-free structural normalization represents a practical and lightweight approach for improving archival compression pipelines without altering or replacing existing compression backends.

**Future Work & Experimental Preview.** To represent a concrete step towards query-ready storage, we have released an **experimental Early Research evolution** of the engine, available in the repository’s `rust_random_access_PREVIEW` directory (and as a pre-compiled Preview release). This preview implements **Independent Row Groups** and **Footer Indexing** to enable  $O(1)$  random access. While the implementation is extremely recent and validated on a non-exhaustive dataset, preliminary benchmarks suggest the viability of this architecture. Initial tests show near-instantaneous retrieval times with a marginal compression trade-off (typically +5% ~ 10%) compared to the Standard version. Furthermore, the updated block-based buffering strategy yielded an **unexpected speedup in full-file decompression** (approx. 20% ~ 35% on mid-sized files), effectively bridging the gap between deep archival efficiency and high-performance query-ready storage.

## Availability

The complete source code, including the reference implementation and experimental artifacts, is available under the MIT License at:

<https://github.com/AndreaLVR/CAST>